

# Generative Models

(Many figures adapted from Stanford CS231n, MIT 6.S191, and Illinois CS 498)

# Outline

Introduction

Variational Autoencoders (VAEs)

- Autoencoders

Generative Adversarial Networks (GANs)

Summary

# Introduction (1/3)

## Supervised vs unsupervised learning

### Supervised Learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn function to map  
 $x \rightarrow y$

**Examples:** Classification, regression, object detection, semantic segmentation, etc.

### Unsupervised Learning

**Data:**  $x$

$x$  is data, no labels!

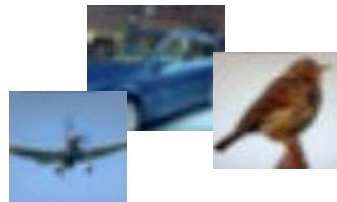
**Goal:** Learn some *hidden* or *underlying structure* of the data

**Examples:** Clustering, feature or dimensionality reduction, etc.

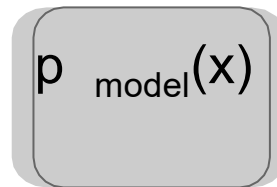
# Introduction (2/3)

- What are **Generative Models**?
  - Generative models are an **Unsupervised Learning** approach
  - Given training data, **generate new samples** from same distribution

Training data  $x \sim p_{\text{data}}$



learning  
→



sampling  
→

Generated samples  $x \sim p_{\text{model}}$



We want to learn  $p_{\text{model}}$  that matches  $p_{\text{data}}$

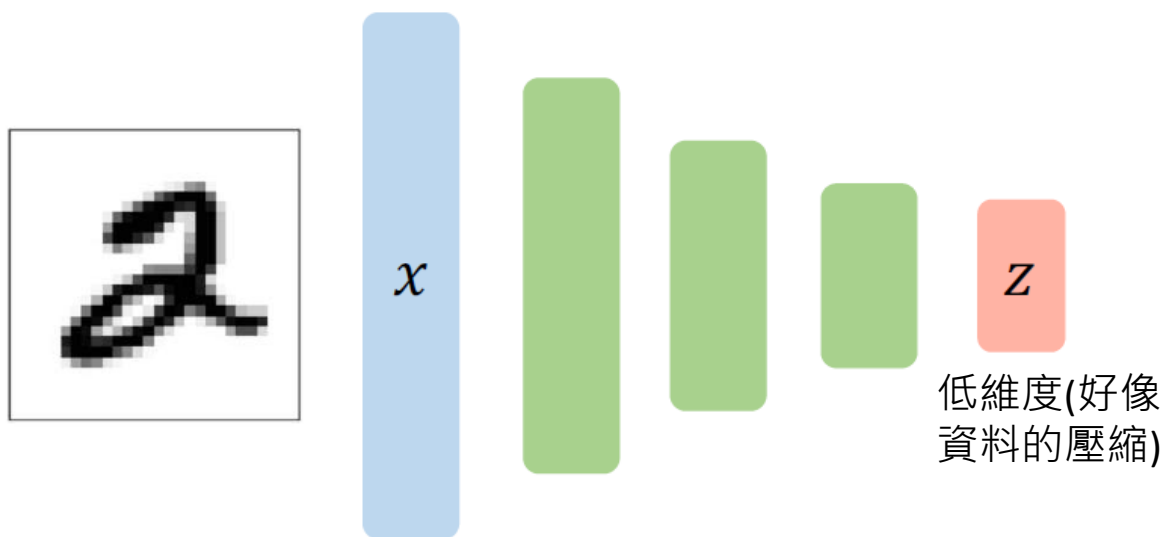
# Introduction (3/3)

- Why generative models?
  - Debiasing
  - Outlier detection
  - and more...
- Introduce two most popular types of generative models today
  - Variational Autoencoders (VAEs)
  - Generative Adversarial Networks (GANs)

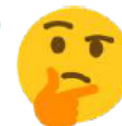
# Autoencoders: Background (1/7)

我沒有告訴他怎麼編碼，  
只是要求他從高維度變成低維度

Unsupervised approach for learning a **lower-dimensional** feature representation from unlabeled training data



Why do we care about a low-dimensional  $z$ ?



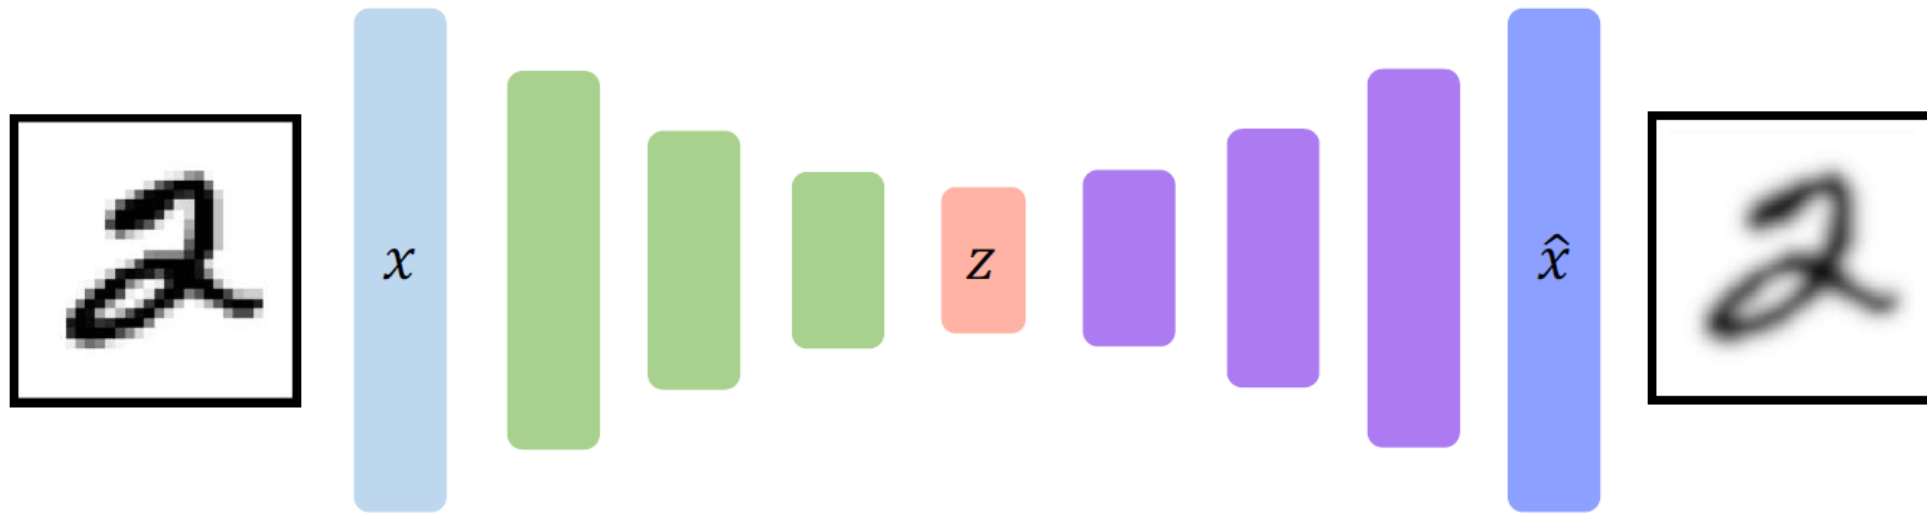
latent 隱藏、潛在

“Encoder” learns mapping from the data,  $x$ , to a low-dimensional latent space,  $z$

# Autoencoders: Background (2/7)

How can we learn this latent space?

Train the model to use these features to **reconstruct the original data**

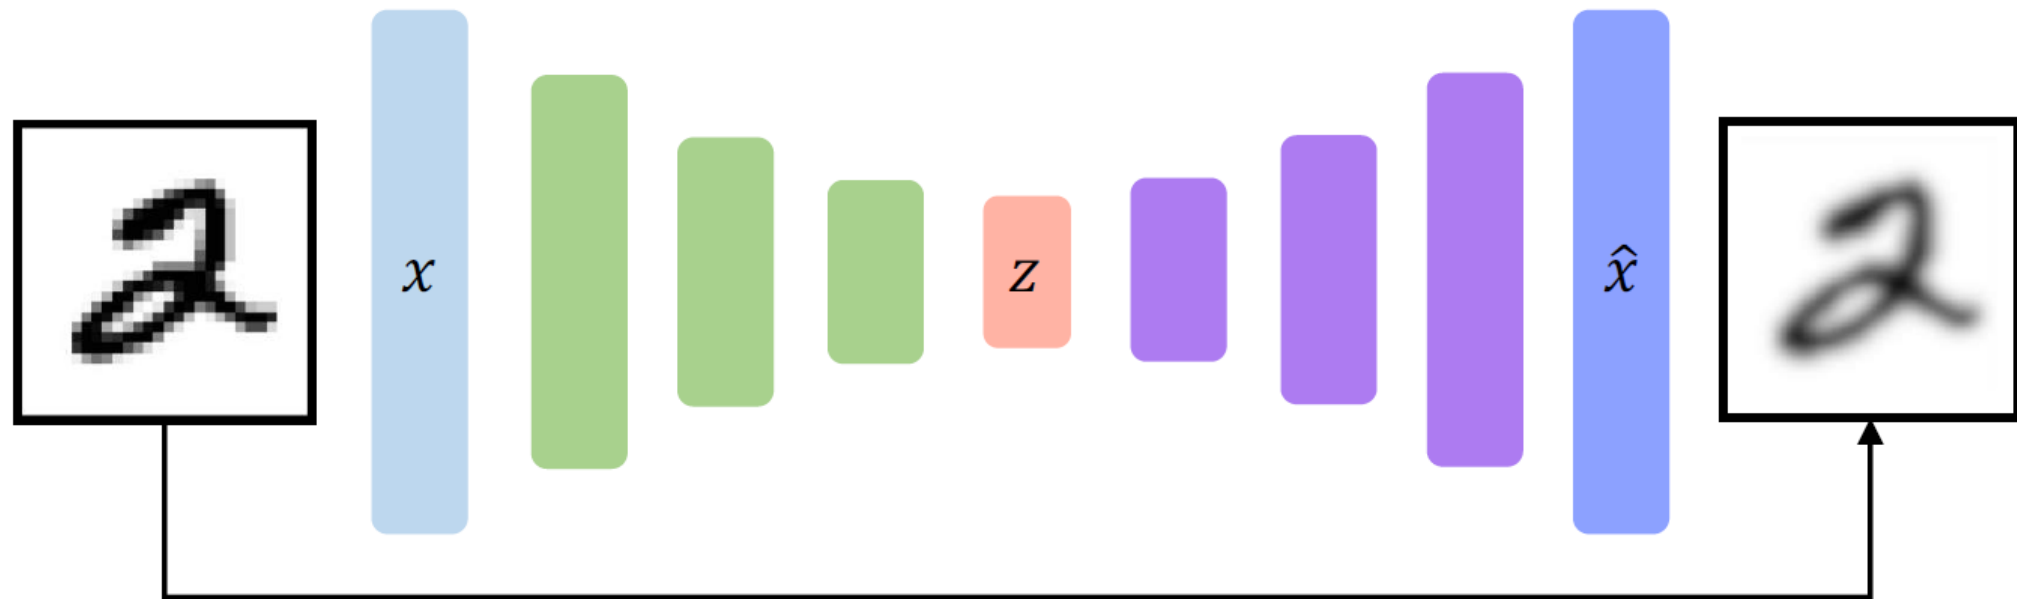


“Decoder” learns mapping back from latent,  $z$ , to a reconstructed observation,  $\hat{x}$

# Autoencoders: Background (3/7)

How can we learn this latent space?

Train the model to use these features to **reconstruct the original data**



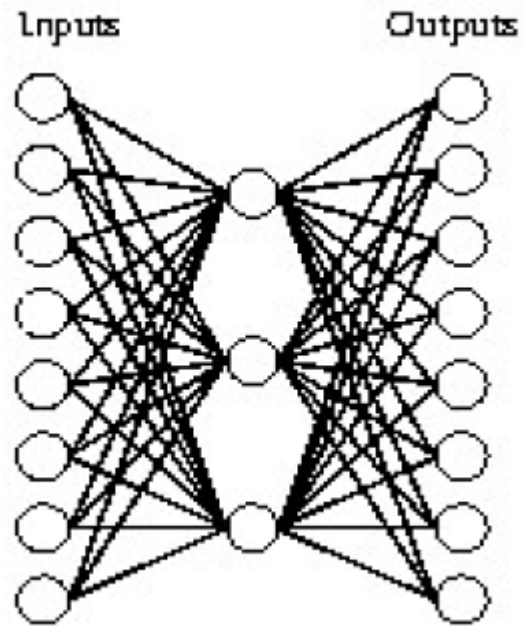
$$\mathcal{L}(x, \hat{x}) = \|x - \hat{x}\|^2$$

Loss function doesn't  
use any labels!!



# Autoencoders: Background (4/7)

**Example 1:** A 8 x 3 x 8 network was trained to learn the identity function



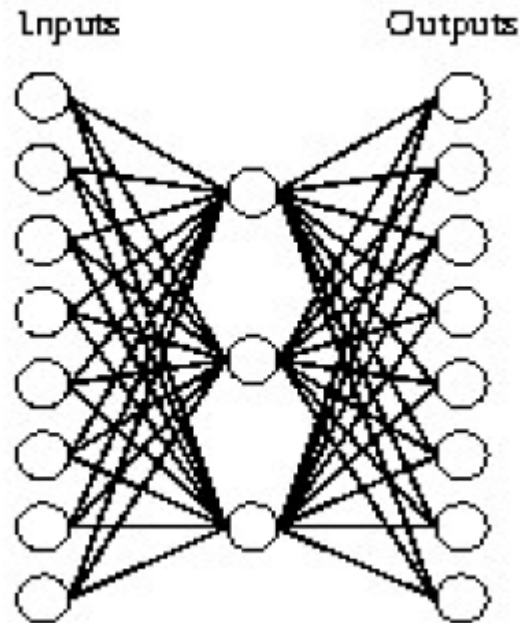
A target function:

Input	Output
10000000	→ 10000000
01000000	→ 01000000
00100000	→ 00100000
00010000	→ 00010000
00001000	→ 00001000
00000100	→ 00000100
00000010	→ 00000010
00000001	→ 00000001

Can this be learned??

# Autoencoders: Background (5/7)

**Example 1:** A 8 x 3 x 8 network was trained to learn the identity function



Input		Hidden Values				Output
10000000	→	.89	.04	.08	→	10000000
01000000	→	.15	.99	.99	→	01000000
00100000	→	.01	.97	.27	→	00100000
00010000	→	.99	.97	.71	→	00010000
00001000	→	.03	.05	.02	→	00001000
00000100	→	.01	.11	.88	→	00000100
00000010	→	.80	.01	.98	→	00000010
00000001	→	.60	.94	.01	→	00000001

Notice that if the encoded values are rounded to **zero** or **one**, the result is the **standard binary encoding** for eight distinct values.

⇒ Eight input features reduce to three, and the **low dimension features** can reconstruct the original data.

# Autoencoders: Background (6/7)

## Example 2: MNIST dataset

Dimensionality of latent space →  
reconstruction quality

Autoencoding is a form of compression!

Smaller latent space will force a larger training bottleneck

2D latent space



5D latent space



Ground Truth



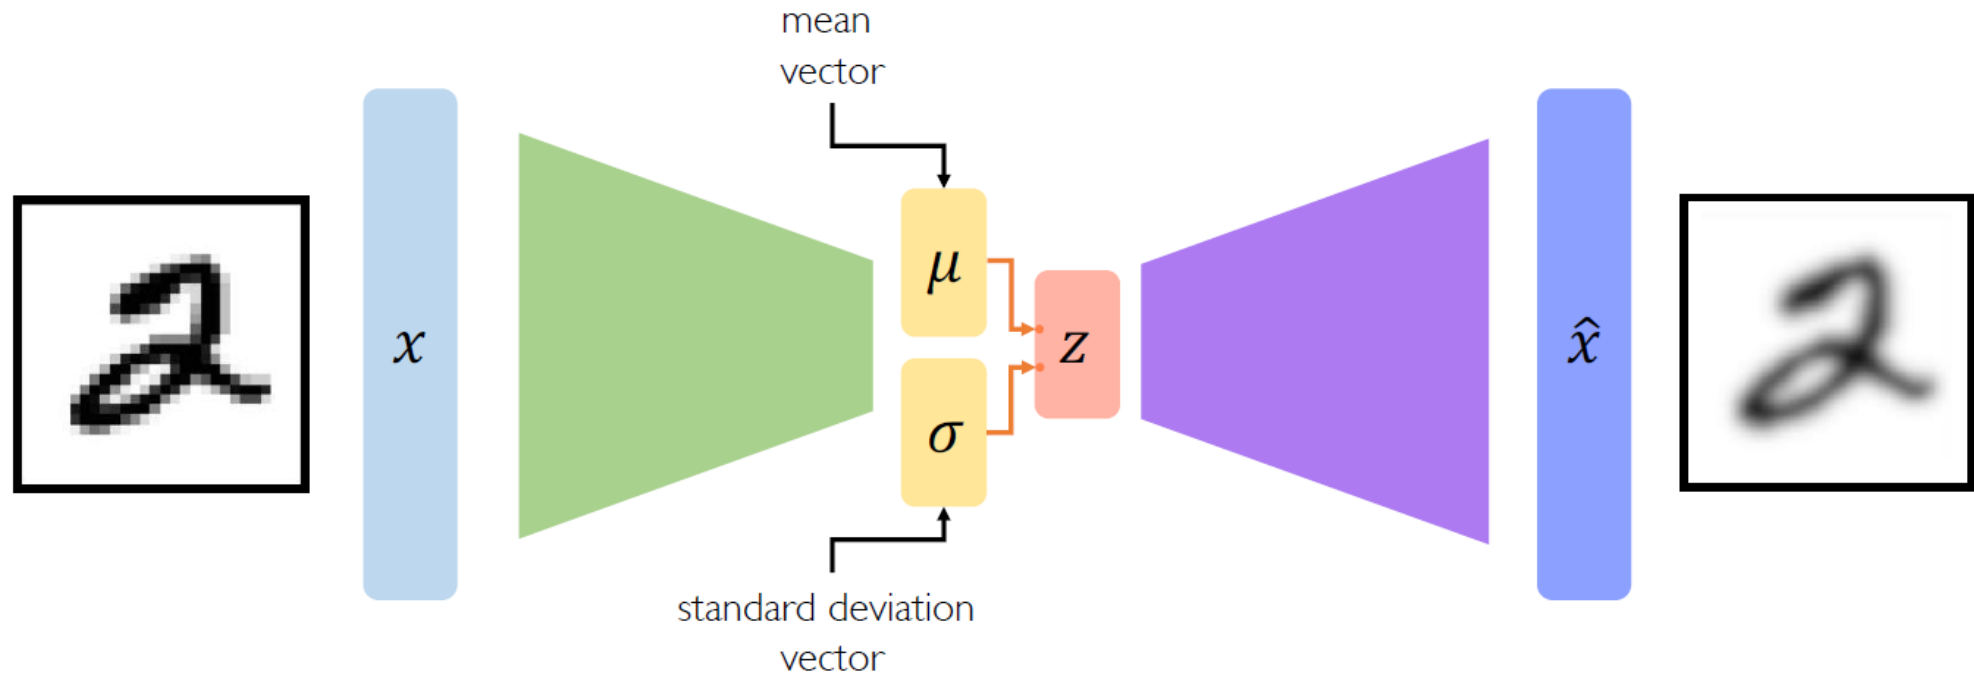
# Autoencoders: Background (7/7)

## Autoencoders Summary

- Autoencoder = Encoder + Decoder (編碼器 + 解碼器)
- Bottleneck hidden layer forces network to learn a compressed latent representation
- Reconstruction loss forces the latent representation to capture (or encode) as much “information” about the data as possible
- Training:
  - Inputs: original input  $X$
  - Targets: original input  $X$  ( $X$  are Not Labels)
- Application: Dimensionality reduction

# Variational Autoencoders (VAEs)

VAEs: **key difference** with traditional autoencoder



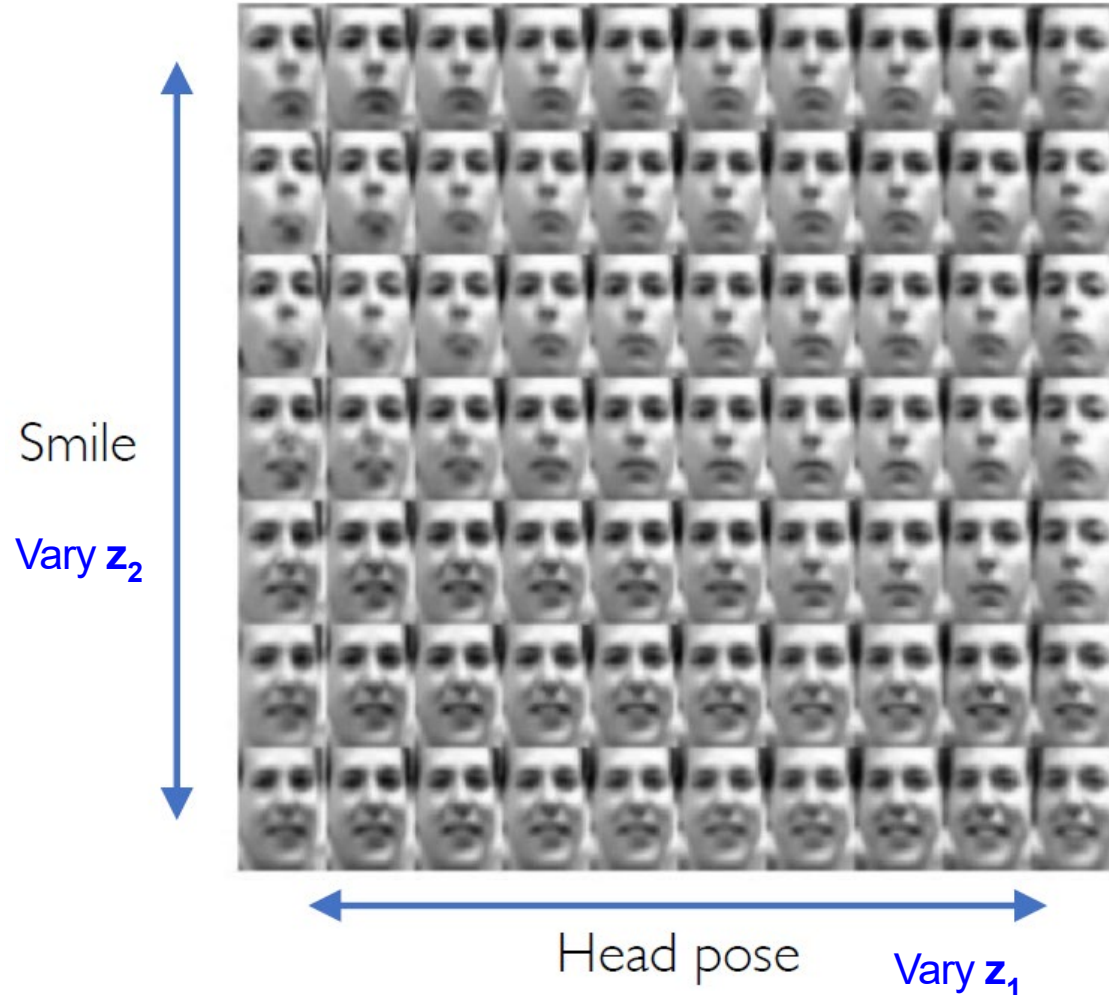
**Variational autoencoders are a probabilistic twist on autoencoders!**

Sample from the mean and standard dev. to compute latent sample

# Variational Autoencoders (VAEs)

Variational Autoencoders:  
Generating Data!

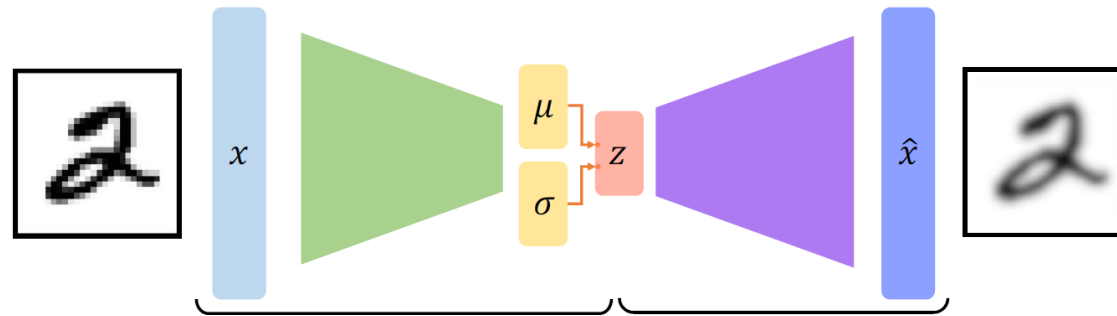
Different dimensions of  $\mathbf{z}$  encode  
interpretable factors of variation



# Variational Autoencoders (VAEs)

## VAEs Summary

- Reparameterization trick to train end-to-end
- Interpret hidden latent variables using perturbation  
⇒ **Generating new examples**



Samples blurrier and lower quality compared to GANs

# Generative Adversarial Networks (GANs) (1/2)

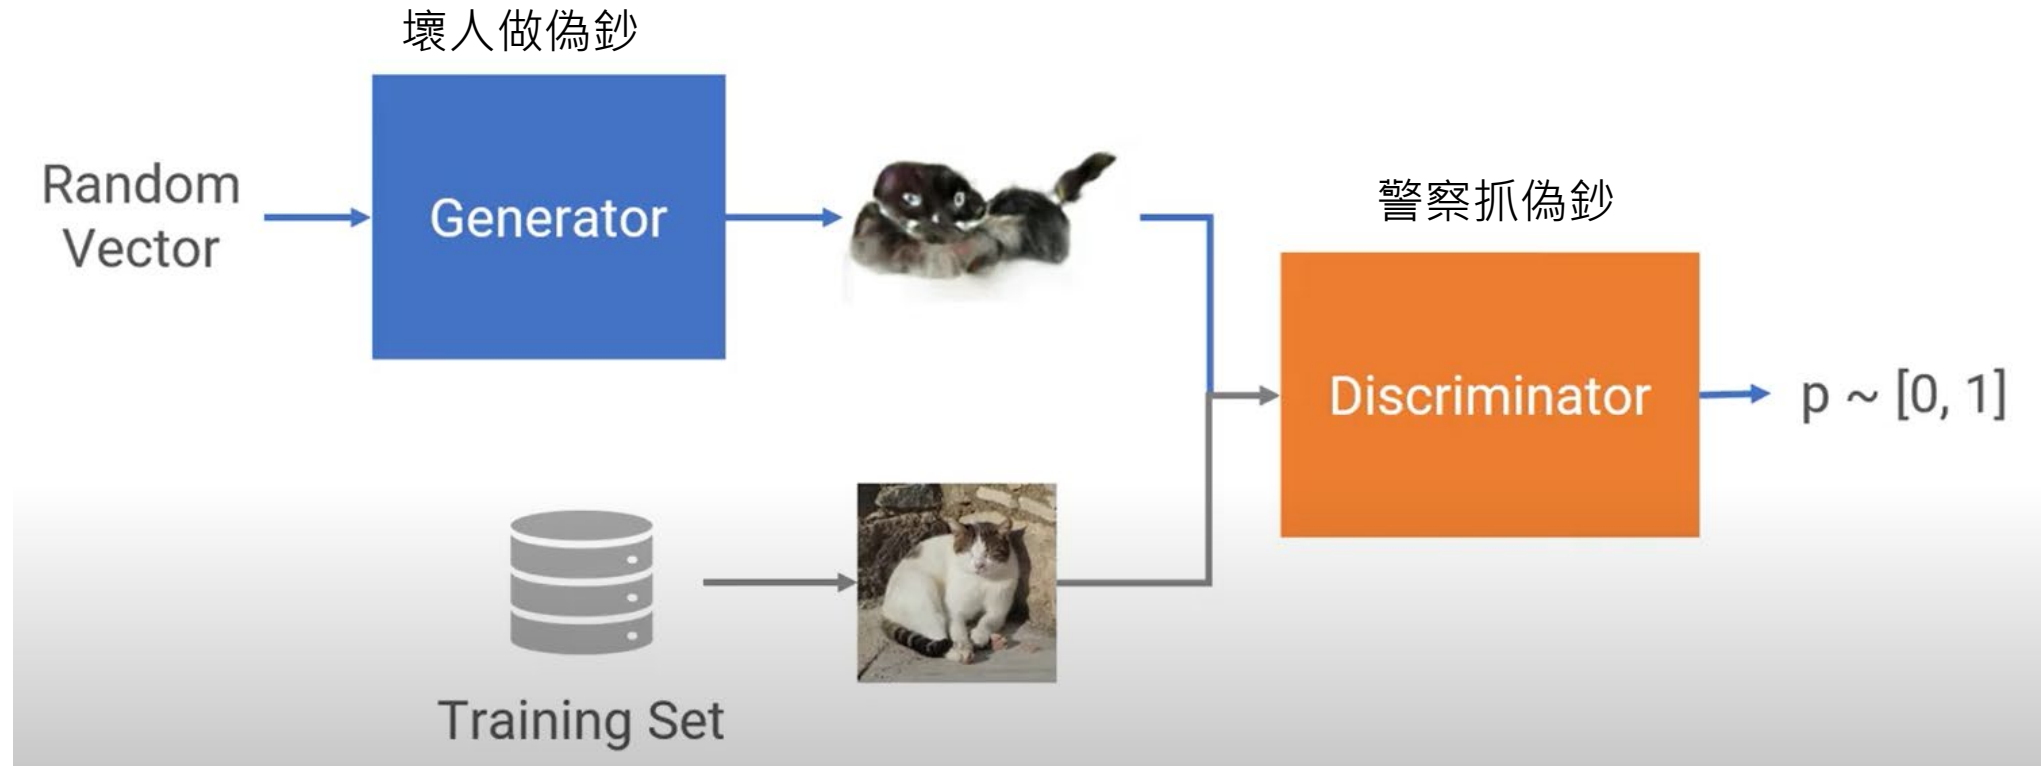
- GANs are an approach to **generative modeling** using **deep learning methods**.
- GANs consist of **two neural networks** that compete against each other during training.
  - The **generator** tries to generate realistic samples that have never been seen before.
  - The **discriminator** tries to identify whether its inputs are real or fake.



# Generative Adversarial Networks (GANs) (2/2)

Train two networks with opposing objectives:

- **Generator:** tries to fool the discriminator by generating real-looking samples
- **Discriminator:** tries to distinguish between generated and real samples



# GAN objective

---

- The discriminator  $D(x)$  should output the probability that the sample  $x$  is real
  - That is, we want  $D(x)$  to be close to 1 for real data and close to 0 for fake
- Expected conditional log likelihood for real and generated data:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

We seed the generator with noise  $z$   
drawn from a simple distribution  $p$

# GAN objective

---

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

- The discriminator wants to correctly distinguish real and fake samples:

$$D^* = \arg \max_D V(G, D)$$

目標：最後讓警察覺得  
真鈔的機率是0.5  
假鈔的機率是0.5

- The generator wants to fool the discriminator:

$$G^* = \arg \min_G V(G, D)$$

- Train the generator and discriminator jointly in *a minimax game*

# GAN Learning Algorithm

---

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

---

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Sample minibatch of  $m$  examples  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  from data generating distribution  $p_{\text{data}}(\mathbf{x})$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(\mathbf{x}^{(i)}) + \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right) \right].$$

**end for**

- Sample minibatch of  $m$  noise samples  $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$  from noise prior  $p_g(\mathbf{z})$ .
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left( 1 - D(G(\mathbf{z}^{(i)})) \right).$$

**end for**

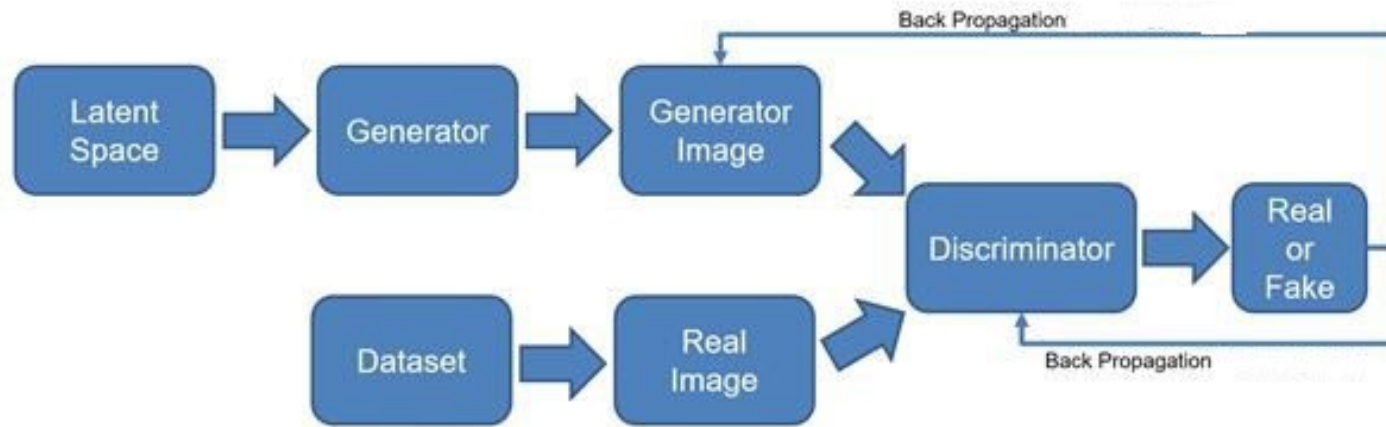
The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

---

# Training of GAN

Repeat the 2 steps:

1. Update the **discriminator network**;
2. Update the **generator network**.



# Phase 1: Update the **Discriminator**

## Train a classifier

1. Generate a batch of **fake samples** by the **generator**;
2. Randomly sample a batch of **real samples**;
3. Inputs:  $\mathbf{X} = [\text{real\_samples}, \text{fake\_samples}]$ ;
4. Targets:  $\mathbf{y} = [\text{True}, \dots, \text{True}, \text{False}, \dots, \text{False}]$ ;
5. Update the **discriminator network** using  $\mathbf{X}$  and  $\mathbf{y}$ .  
(**freeze generator's** parameters)

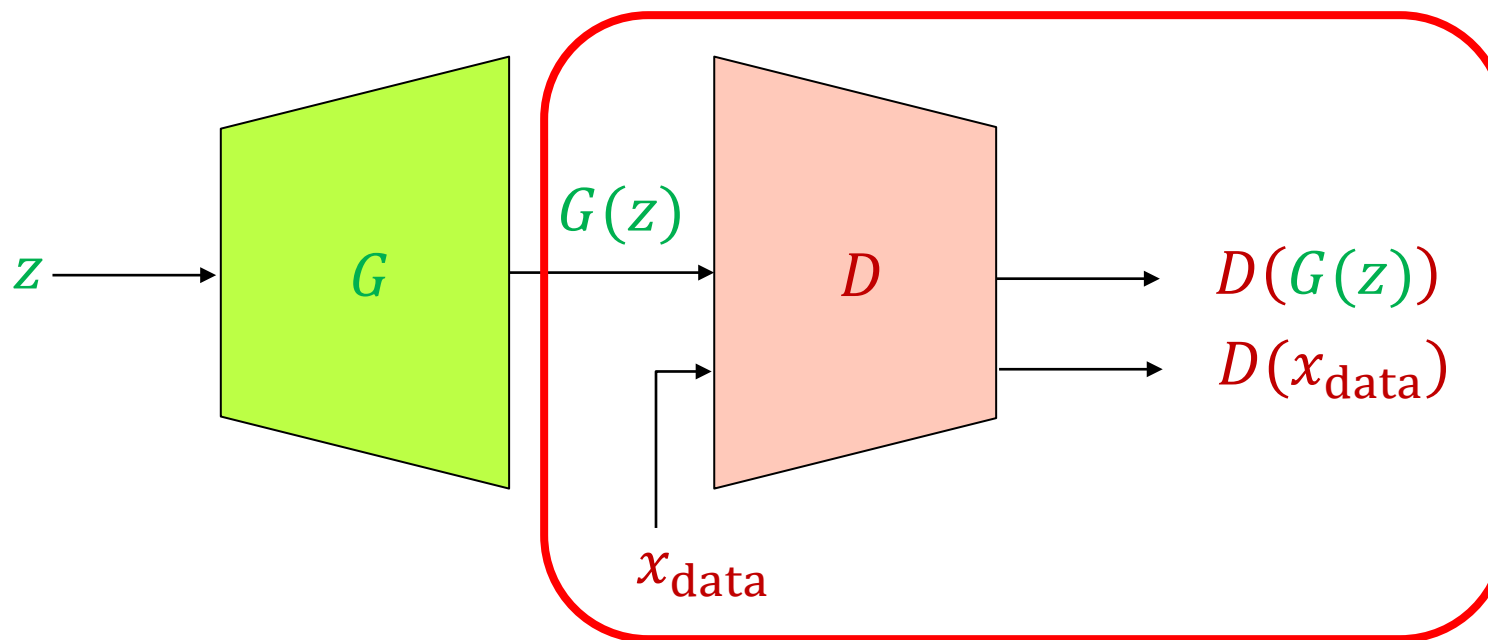
# Phase 2: Update the **Generator**

Connect the **generator** and **discriminator**  
(**freeze discriminator's** parameters).

# GAN: Conceptual picture

---

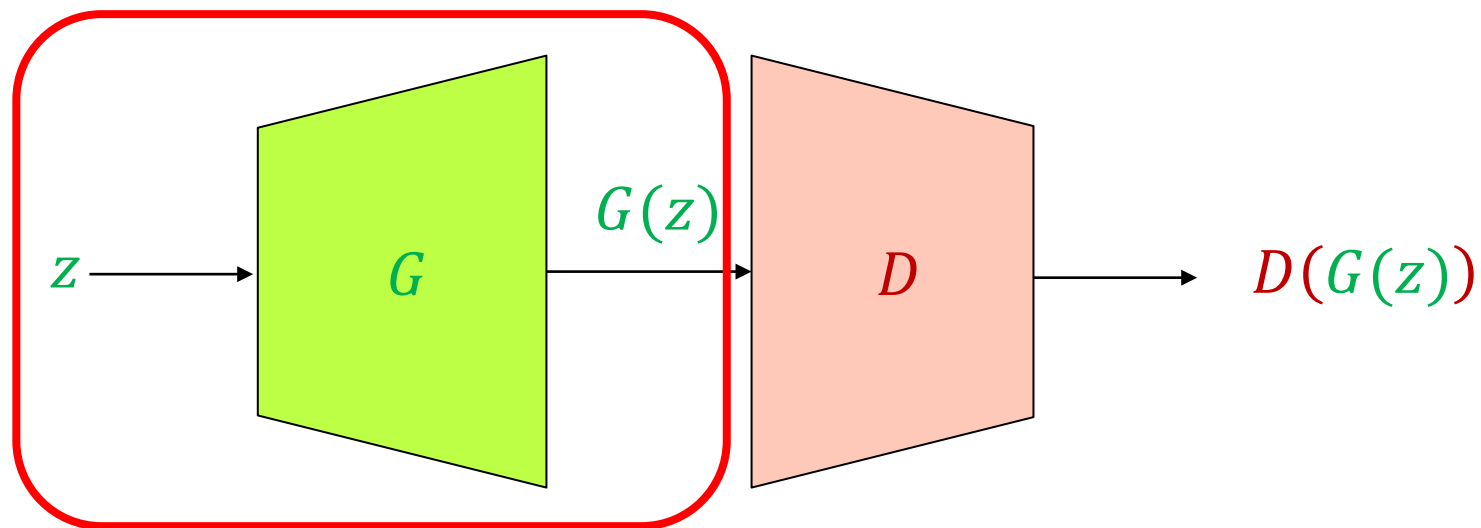
- Update discriminator
  - push  $D(x_{\text{data}})$  close to 1 and  $D(G(z))$  close to 0
  - freeze generator's parameters



# GAN: Conceptual picture

---

- Update generator: increase  $D(G(z))$ 
  - Requires back-propagating through the composed generator-discriminator network
  - Freezes discriminator's parameter

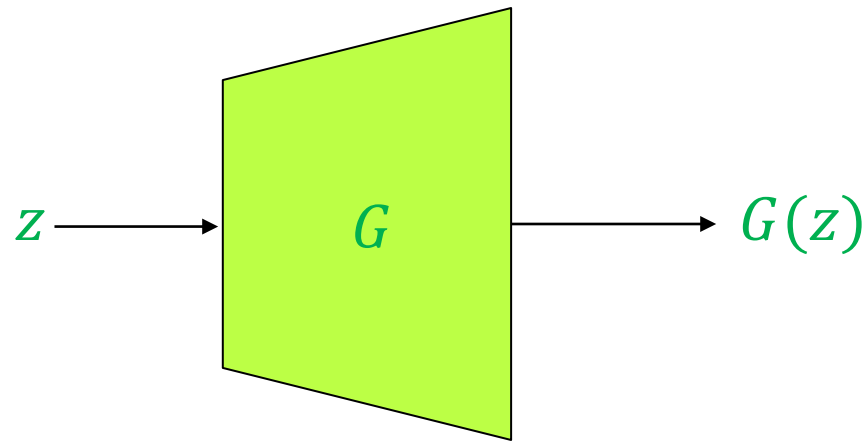




# GAN: Conceptual picture

---

- Test time – the discriminator is discarded



# 2017: Explosion of GANs

See also: <https://github.com/soumith/ganhacks> for tips and tricks for trainings GANs

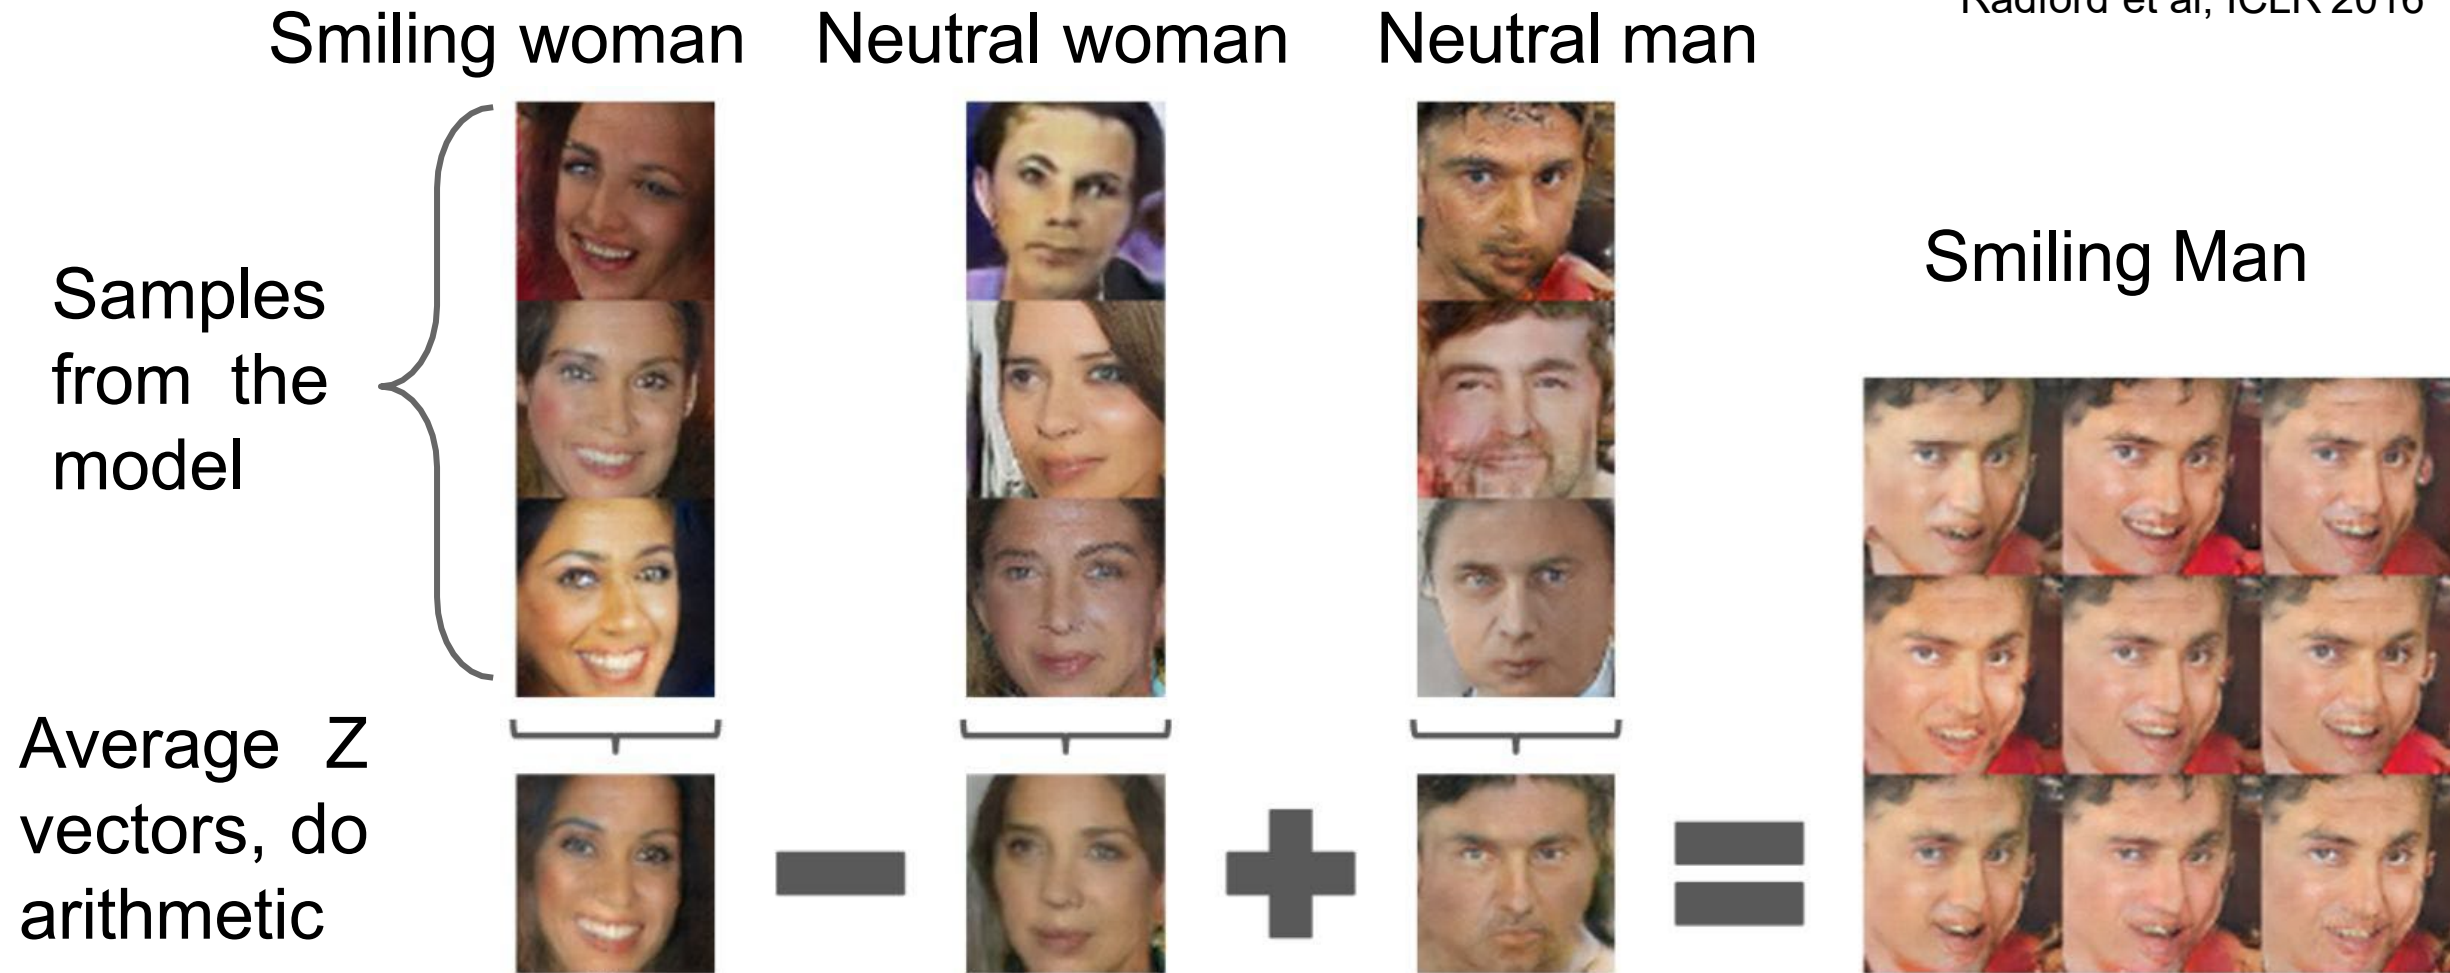
## “The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>

# Generative Adversarial Nets: Interpretable Vector Math

Radford et al, ICLR 2016





# Generative Adversarial Nets: Interpretable Vector Math

Glasses man



No glasses man



No glasses woman



Radford et al,  
ICLR 2016

Woman with glasses



−

+

=

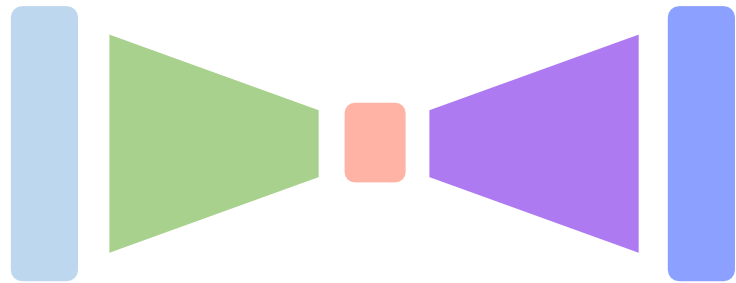
# Which face is fake?



# Generative Models: Summary

## Autoencoders and Variational Autoencoders (VAEs)

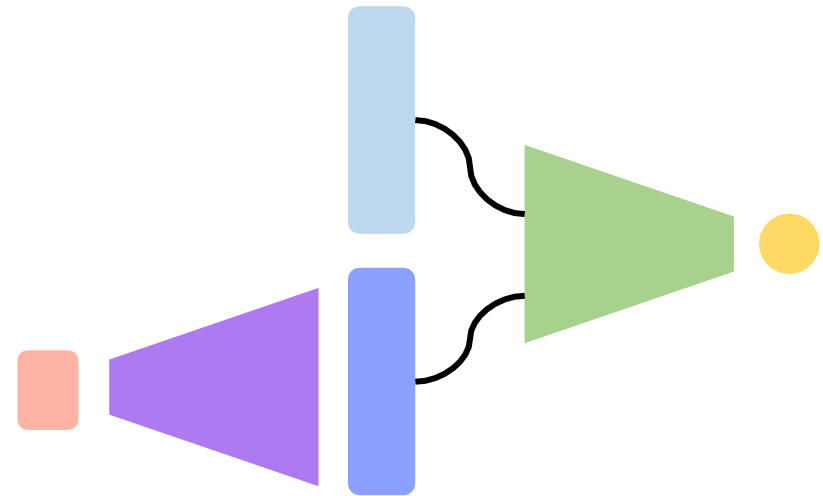
Learn lower-dimensional latent space and sample to generate input reconstructions



VAEs: **Explicit** density method

## Generative Adversarial Networks (GANs)

Competing generator and discriminator networks



GANs: **Implicit** density method

# Resources

Stanford CS231n: Convolutional Neural Networks for Visual Recognition

MIT 6.S191: Introduction to Deep Learning

Illinois CS 498: Introduction to Deep Learning